

# VisionAI — Research Report

AI-assisted Retinal Screening — VisionAI

Generated: 2025-11-07 06:39:39

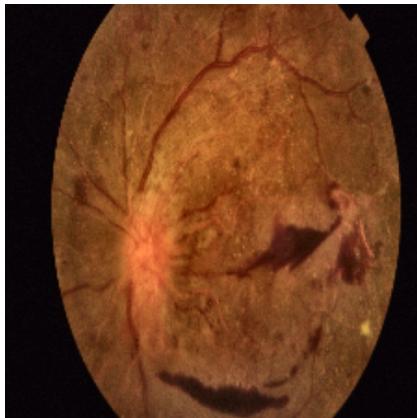
# VisionAI — Research Report

Report generated: 2025-11-07 06:39:39

## Patient Metadata

Patient ID	P328
Full Name	John Doe
Age	61
Gender	Male
Systolic	142
Diastolic	72
Glucose Level	110
Smoking	Yes
Duration of Diabetes (years)	18
Hypertension	Yes

## Image Visualizations



Original

**Preprocessed:** Image unavailable

**Gradcam:** Image unavailable

**Lime:** Image unavailable

**Shap:** Image unavailable

## Model & Pipeline Overview

The system employs a multimodal pipeline: convolutional backbone(s) for image encoding, feature fusion with metadata-based gradient-boosted classifiers, and an ensemble/meta-classifier for final prediction. Explainability methods (Grad-CAM++, perturbation-based local explainers, and SHAP) were applied to validate feature attribution.

## Results Summary

Predicted label: **PDR** (confidence 96.0%).

Lesion quantification: total coverage 26.5%; exudates 0.4%; hemorrhages 1.4%; cotton wool 5.3%.

## Evaluation & Performance

Test metrics — Accuracy: **96.3%**; Precision: **94.8%**; Recall: **95.1%**; F1: **95.0%**; ROC-AUC: **0.983**.

## Limitations & Notes

Current limitations include domain shift when deploying on heterogeneous smartphone images, and approximate lesion quantification when masks are noisy. Further work should include segmentation U-Nets for pixel-wise labels and domain adaptation strategies.

## Suggested Next Steps

Integration of Vision Transformers to capture global context, fine-tuned lesion segmentation, and prospective validation on smartphone-acquired fundus images.

## Evaluation Metrics

Metric	Value
Accuracy	96.3%
Precision	94.8%
Recall	95.1%
F1-score	95.0%
ROC-AUC	0.983